

CS 236: Deep Generative Models

Stefano Ermon and Aditya Grover
Stanford University

URL: deepgenerativemodels.github.io

Introduction

Challenge: understand complex, unstructured inputs



Computer Vision



Natural Language Processing



Computational Speech



Robotics

Introduction



Richard Feynman: “*What I cannot create, I do not understand*”

Generative modeling: “*What I understand, I can **create***”

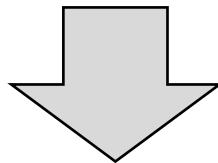
Generative Modeling: Computer Graphics

How to generate natural images with a computer?

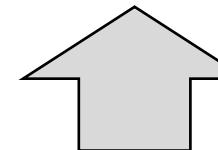
High level
description

Cube(color=blue, position=, size=, ...)
Cylinder(color=red, position=, size=,..)

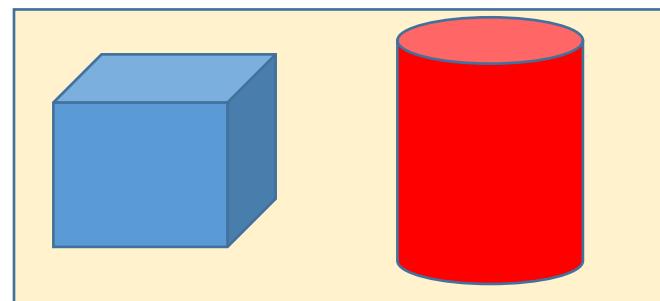
Generation (graphics)



Inference (vision)



Raw sensory
outputs



Our models will have **similar structure (generation + inference)**

Statistical Generative Models

Statistical generative models are **learned from data**



Data
(e.g., images of bedrooms)

+



Prior Knowledge
(e.g., physics, materials, ..)

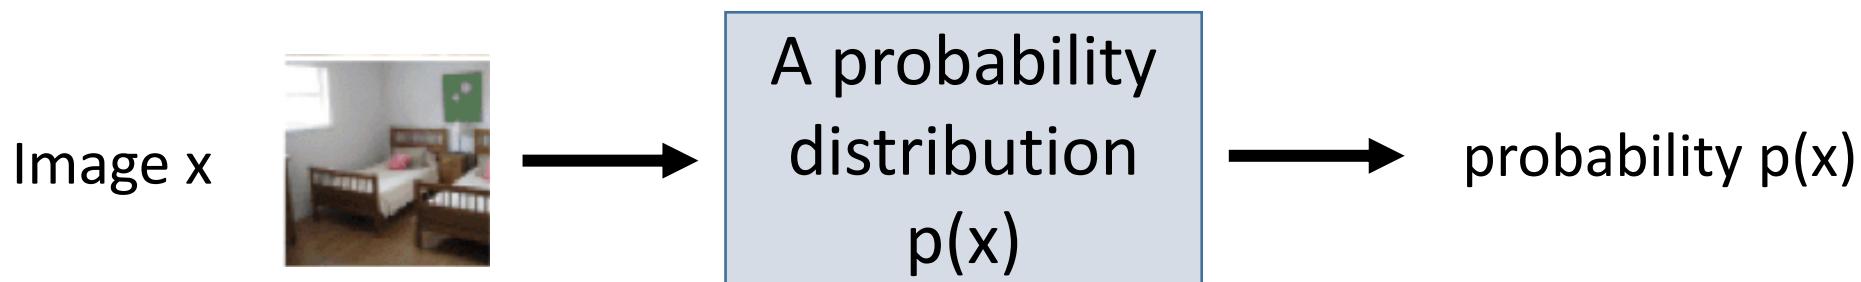
Priors are always necessary, but there is a spectrum



Statistical Generative Models

A statistical generative model is a **probability distribution** $p(x)$

- **Data:** samples (e.g., images of bedrooms)
- **Prior knowledge:** parametric form (e.g., Gaussian?), loss function (e.g., maximum likelihood?), optimization algorithm, etc.

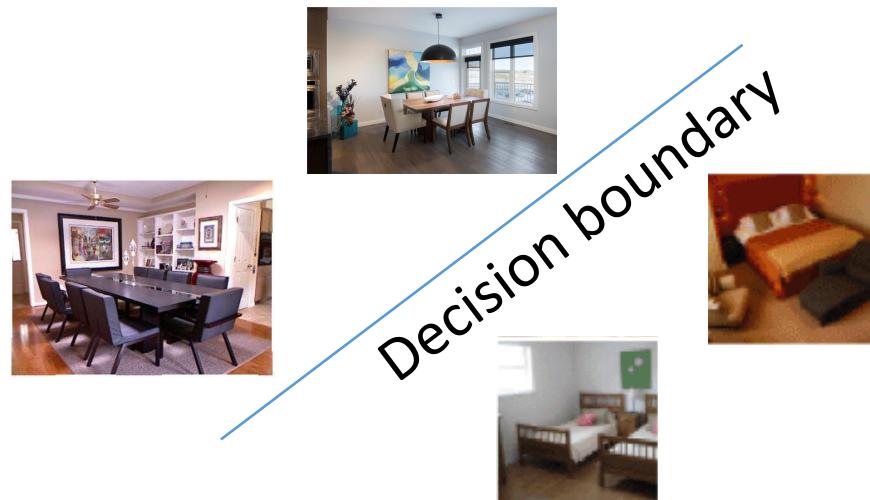


It is generative because **sampling from $p(x)$ generates new images**

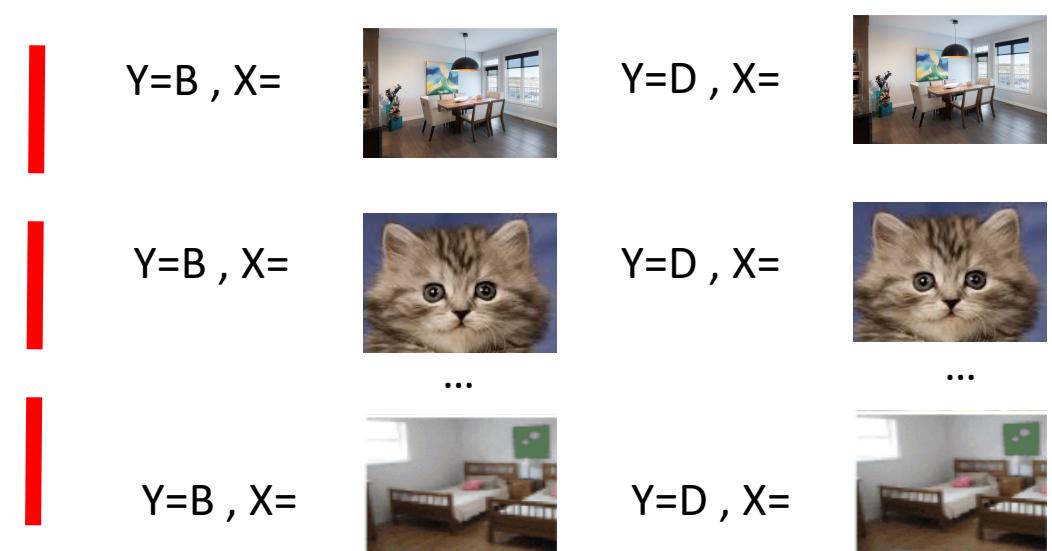


Discriminative vs. generative

Discriminative: classify bedroom vs. dining room



Generative: generate X



The input image X is always given. **Goal:** a good decision boundary, via **conditional distribution**

$$P(Y = \text{Bedroom} \mid X=)$$

Ex: logistic regression, convolutional net, etc.

The input X is **not** given. Requires a model of the **joint distribution**

$$P(Y = \text{Bedroom}, X=)$$

Discriminative vs. generative

Joint and conditional are related via **Bayes Rule**:

$P(Y = \text{Bedroom} | X=$



$P(Y = \text{Bedroom}, X=$



$P(X=$



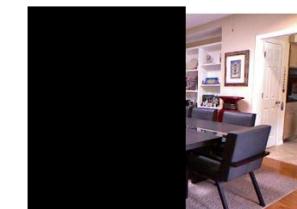
Discriminative: X is always given, does not need to model

Therefore it cannot handle missing data

$P(X=$



$P(Y = \text{Bedroom} | X=$



Conditional Generative Models

Class **conditional generative models** are also possible:

$$P(X = \text{Bedroom})$$



It's often useful to condition on rich side information Y

$$P(X = \text{Bedroom} | \text{option} = \text{"A black table with 6 chairs"})$$



A discriminative model is a very simple conditional generative model of Y:

$$P(Y = \text{Bedroom} | X = \text{Bedroom})$$



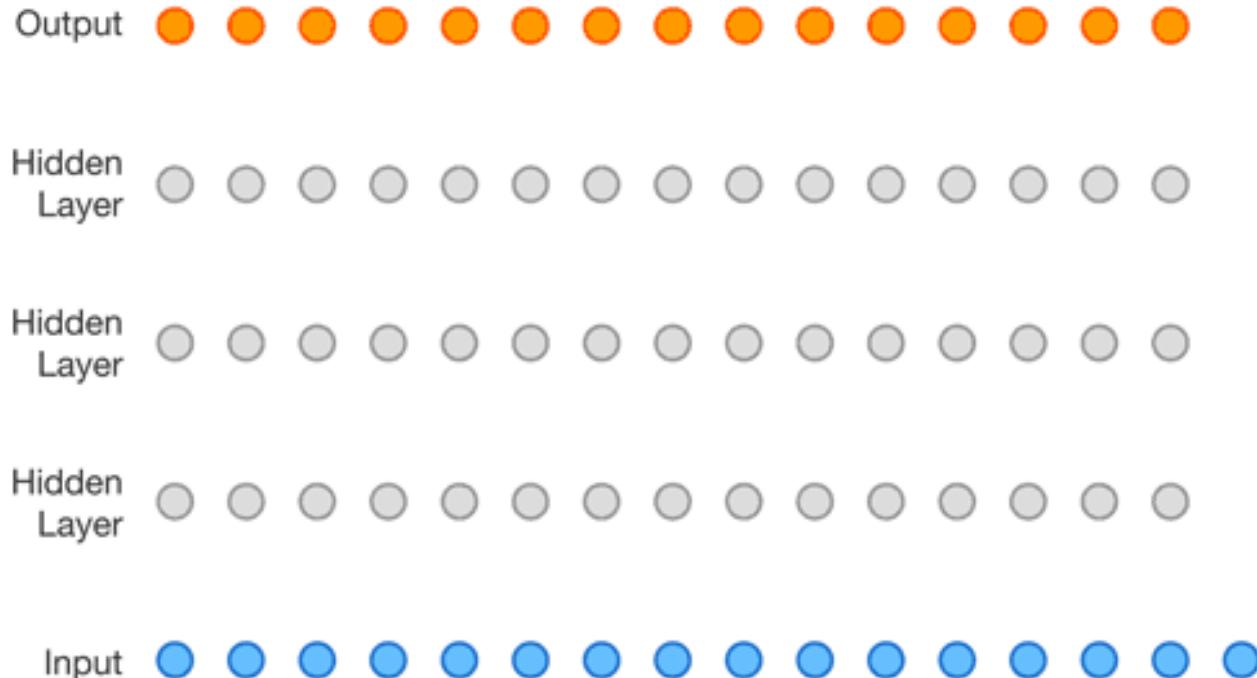
Progressive Growing of GANs



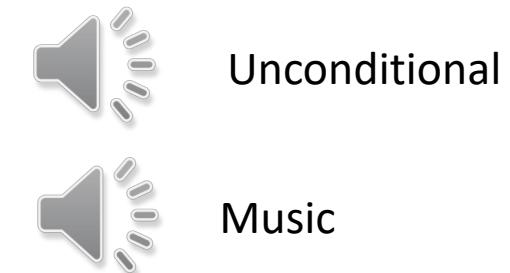
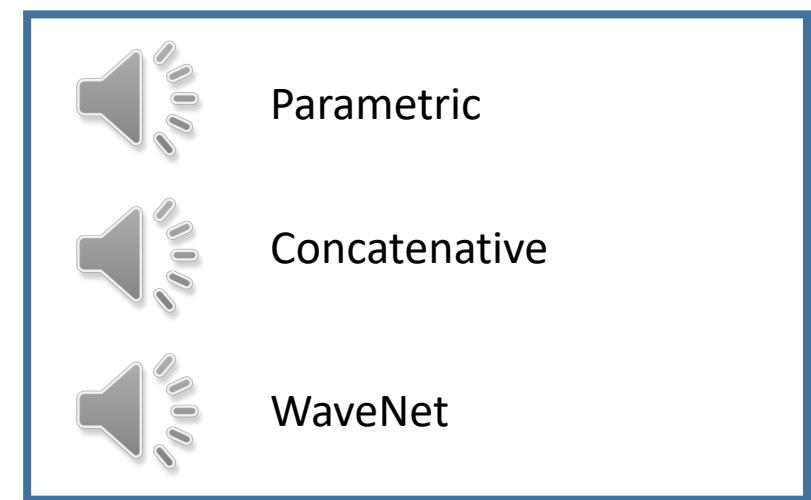
Karras et al., 2018

WaveNet

Generative model of speech signals



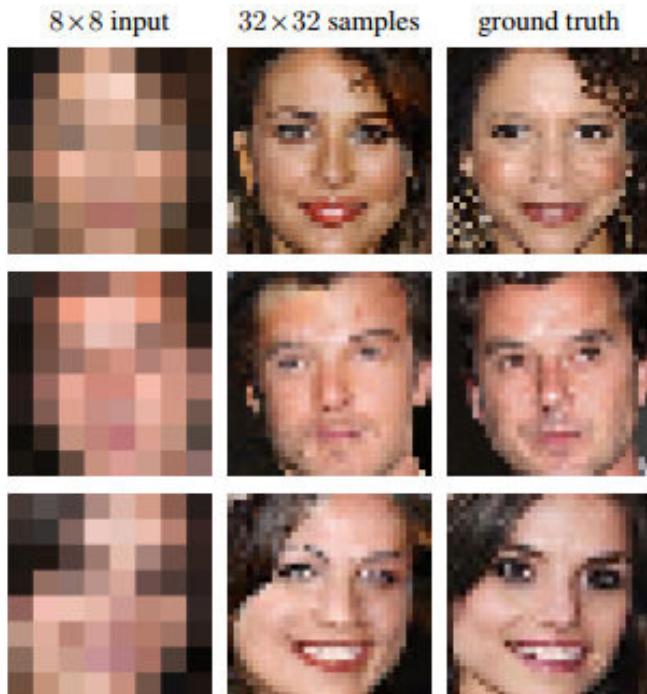
Text to Speech



van den Oord et al, 2016c

Image Super Resolution

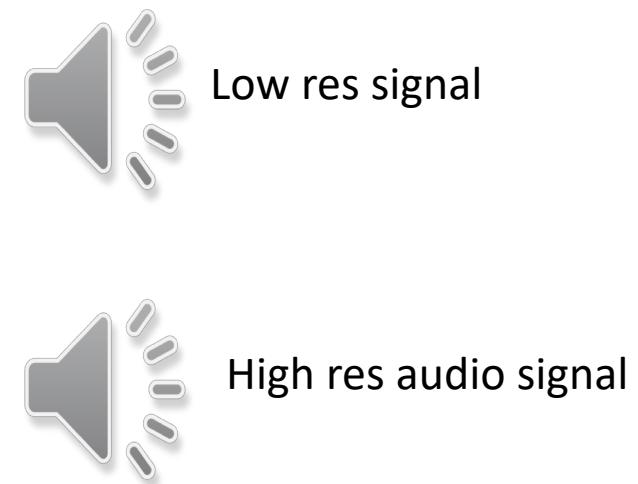
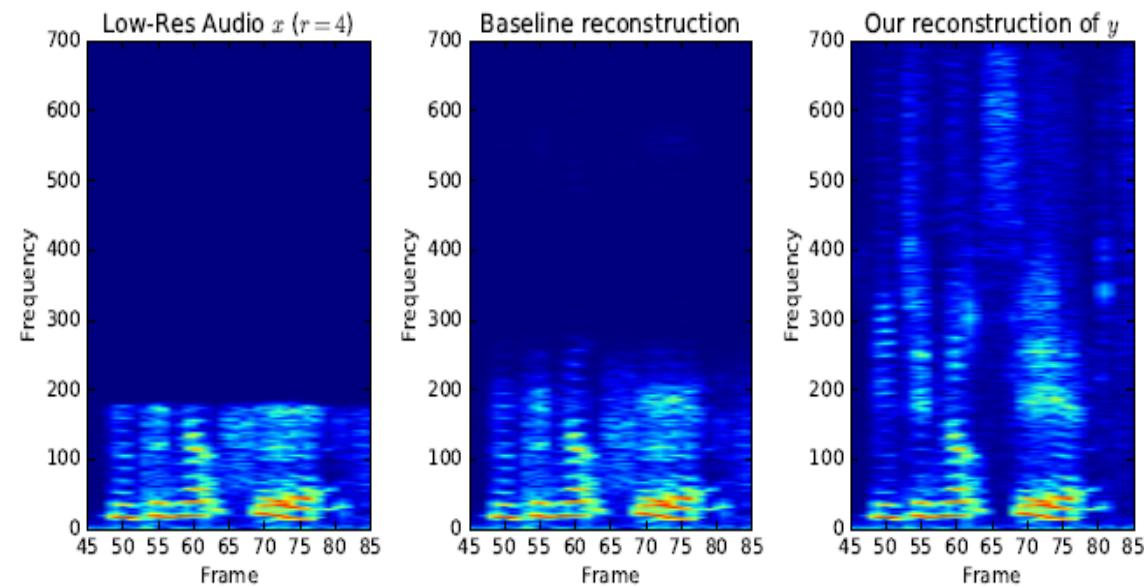
Conditional generative model $P(\text{high res image} \mid \text{low res image})$



Ledig et al., 2017

Audio Super Resolution

Conditional generative model $P(\text{high-res signal} \mid \text{low-res audio signal})$



Kuleshov et al., 2017

Machine Translation

Conditional generative model $P(\text{ English text} | \text{ Chinese text})$

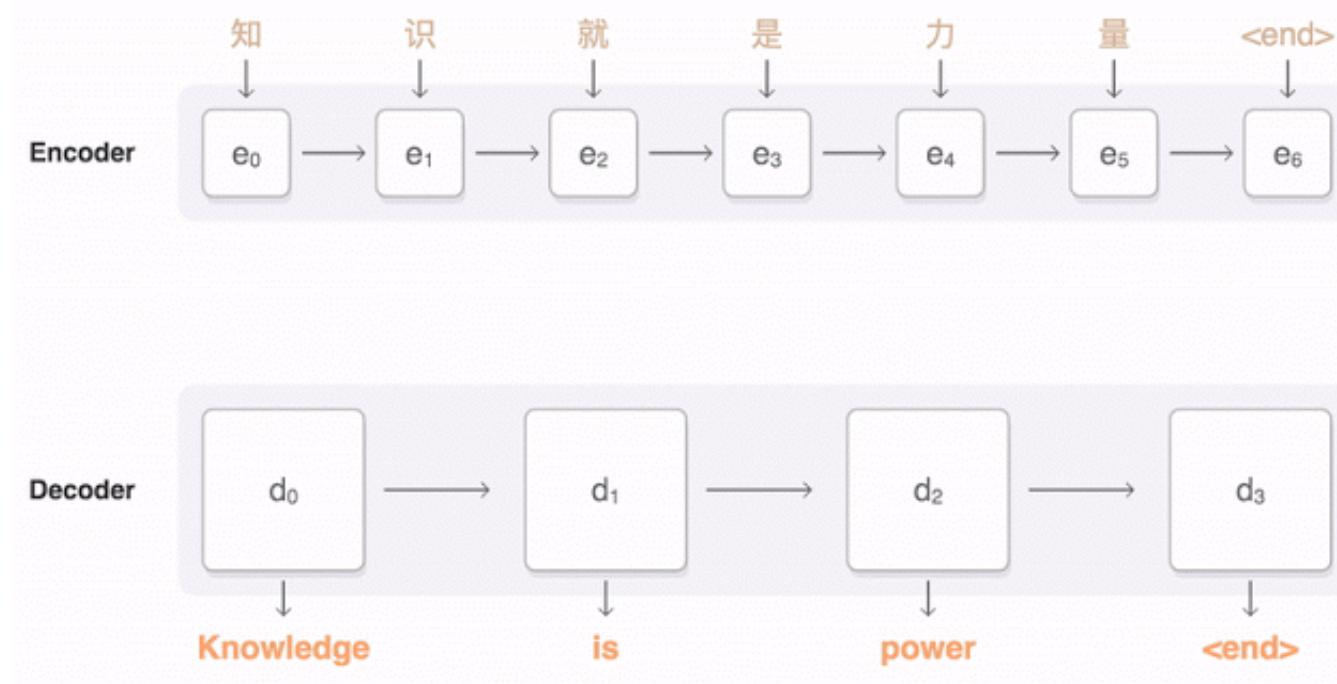


Figure from Google AI research blog.

Image Translation

Conditional generative model $P(\text{ zebra images} | \text{ horse images})$



Zhu et al., 2017

Imitation Learning

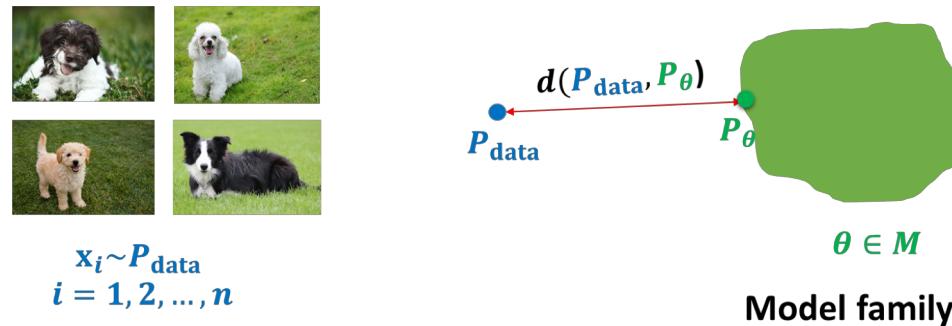
Conditional generative model $P(\text{actions} \mid \text{past observations})$



Li et al., 2017

Roadmap and Key Challenges

- **Representation:** how do we model the joint distribution of many random variables?
 - Need compact representation
- **Learning:** what is the right way to compare probability distributions?



- **Inference:** how do we invert the generation process (e.g., vision as inverse graphics)?
 - Unsupervised learning: recover high-level descriptions (features) from raw data

Syllabus

- Fully observed likelihood-based models
 - Autoregressive
 - Flow-based models
- Latent variable models
 - Variational learning
 - Inference amortization
 - Variational autoencoder
- Implicit generative models
 - Two sample tests, embeddings, F-divergences
 - Generative Adversarial Networks
- Learn about algorithms, theory & applications

Prerequisites

- Basic knowledge about machine learning from at least one of CS 221, 228, 229 or 230.
- Basic knowledge of probabilities and calculus:
 - Gradients, gradient-descent optimization, backpropagation
 - Random variables, independence, conditional independence
 - Bayes rule, chain rule, change of variables formulas
- Proficiency in some programming language, preferably Python, required.

Logistics

- Class webpage: <https://deepgenerativemodels.github.io/>
- <http://piazza.com/stanford/fall2018/cs236>
- There is no required textbook. Reading materials and course notes will be provided.
- Suggested Reading: *Deep Learning* by Ian Goodfellow, Yoshua Bengio, Aaron Courville. Online version available free [here](#).
- Lecture notes are under construction
- Teaching Assistants: Yang Song, Jiaming Song, Rui Shu, Casey Chu, Nish Khandwala
- Office hours: See calendar on class website

Logistics – Grading policies

- Grading Policy:
 - Three homeworks (15% each): mix of conceptual and programming based questions
 - Midterm: 15%
 - Course Project: 40%
 - Proposal: 5%
 - Progress Report: 10%
 - Poster Presentation: 10%
 - Final Report: 15%

Projects

- Course projects will be done in groups of up to 3 students and can fall into one or more of the following categories:
 - Application of deep generative models on a novel task/dataset
 - Algorithmic improvements into the evaluation, learning and/or inference of deep generative models
 - Theoretical analysis of any aspect of existing deep generative models
- Teaching staff will suggest possible projects
- We will provide Google Cloud coupons